

In [1]: *# Assignment 4: ECG Anomaly Detection using Autoencoders*

```
import numpy as np
import pandas as pd
import tensorflow as tf
from tensorflow import keras
from keras.models import Model
from keras.layers import Input, Dense
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
import matplotlib.pyplot as plt

# =====
# 1. Load ECG Dataset
# =====
url = "http://storage.googleapis.com/download.tensorflow.org/data/ecg.csv"
data = pd.read_csv(url, header=None)

# Last column = Labels (0 = anomaly, 1 = normal)
X = data.iloc[:, :-1].values
y = data.iloc[:, -1].values

# =====
# 2. Preprocessing
# =====
scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(X)

# Train only on "normal" samples (Label = 1)
X_normal = X_scaled[y == 1]

# Split into train-test sets
x_train, x_test = train_test_split(X_normal, test_size=0.2, random_state=42)

# Also keep anomalies for evaluation
x_test_all = X_scaled
y_test_all = y

# =====
# 3. Build Autoencoder
# =====
input_dim = X.shape[1] # number of features

input_layer = Input(shape=(input_dim,))
# Encoder
encoded = Dense(64, activation="relu")(input_layer)
encoded = Dense(32, activation="relu")(encoded)
encoded = Dense(16, activation="relu")(encoded)
# Decoder
decoded = Dense(32, activation="relu")(encoded)
decoded = Dense(64, activation="relu")(decoded)
decoded = Dense(input_dim, activation="sigmoid")(decoded)

autoencoder = Model(input_layer, decoded)
autoencoder.compile(optimizer="adam", loss="mse")

# =====
# 4. Train Autoencoder
```

```

# =====
history = autoencoder.fit(x_train, x_train,
                          epochs=20,
                          batch_size=512,
                          validation_data=(x_test, x_test),
                          verbose=1)

# =====
# 5. Plot Training Loss
# =====
plt.plot(history.history['loss'], label="Train Loss")
plt.plot(history.history['val_loss'], label="Validation Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.title("Autoencoder Training Loss")
plt.show()

# =====
# 6. Reconstruction Error
# =====
reconstructions = autoencoder.predict(x_test_all)
mse = np.mean(np.power(x_test_all - reconstructions, 2), axis=1)

# Set threshold = mean + std of training reconstruction error
threshold = np.mean(mse) + np.std(mse)
print("Reconstruction error threshold:", threshold)





















# =====
# 7. Anomaly Detection
# =====
y_pred = [0 if e > threshold else 1 for e in mse]


from sklearn.metrics import classification_report, confusion_matrix

print("\nConfusion Matrix:")
print(confusion_matrix(y_test_all, y_pred))

print("\nClassification Report:")
print(classification_report(y_test_all, y_pred))

```

Epoch 1/20
5/5  1s 51ms/step - loss: 0.0281 - val_loss: 0.0270
Epoch 2/20
5/5  0s 26ms/step - loss: 0.0266 - val_loss: 0.0255
Epoch 3/20
5/5  0s 25ms/step - loss: 0.0248 - val_loss: 0.0226
Epoch 4/20
5/5  0s 29ms/step - loss: 0.0217 - val_loss: 0.0185
Epoch 5/20
5/5  0s 28ms/step - loss: 0.0175 - val_loss: 0.0146
Epoch 6/20
5/5  0s 26ms/step - loss: 0.0136 - val_loss: 0.0113
Epoch 7/20
5/5  0s 25ms/step - loss: 0.0104 - val_loss: 0.0086
Epoch 8/20
5/5  0s 25ms/step - loss: 0.0081 - val_loss: 0.0069
Epoch 9/20
5/5  0s 24ms/step - loss: 0.0064 - val_loss: 0.0060
Epoch 10/20
5/5  0s 26ms/step - loss: 0.0056 - val_loss: 0.0055
Epoch 11/20
5/5  0s 23ms/step - loss: 0.0053 - val_loss: 0.0052
Epoch 12/20
5/5  0s 24ms/step - loss: 0.0049 - val_loss: 0.0051
Epoch 13/20
5/5  0s 24ms/step - loss: 0.0049 - val_loss: 0.0050
Epoch 14/20
5/5  0s 26ms/step - loss: 0.0048 - val_loss: 0.0049
Epoch 15/20
5/5  0s 24ms/step - loss: 0.0046 - val_loss: 0.0048
Epoch 16/20
5/5  0s 22ms/step - loss: 0.0046 - val_loss: 0.0048
Epoch 17/20
5/5  0s 24ms/step - loss: 0.0047 - val_loss: 0.0048
Epoch 18/20
5/5  0s 23ms/step - loss: 0.0046 - val_loss: 0.0047
Epoch 19/20
5/5  0s 24ms/step - loss: 0.0045 - val_loss: 0.0047
Epoch 20/20
5/5  0s 27ms/step - loss: 0.0045 - val_loss: 0.0047



157/157 ————— 0s 1ms/step
Reconstruction error threshold: 0.02259892845003093

Confusion Matrix:
[[874 1205]
[92 2827]]

Classification Report:

	precision	recall	f1-score	support
0.0	0.90	0.42	0.57	2079
1.0	0.70	0.97	0.81	2919
accuracy			0.74	4998
macro avg	0.80	0.69	0.69	4998
weighted avg	0.79	0.74	0.71	4998

In []: